

● AVAILABLE FOR SENIOR CONTRACTS · REMOTE · GLOBAL

Jorge Aguirre Romero

Senior **Cloud & AI Platform** Architect — GPU & Agentic AI for
Regulated Enterprises

📍 Madrid, Spain · Remote · Open to relocation (EU / US) 📞 +34 687 172 900

✉ aguirre_coslada@hotmail.com in/jorge-aguirre-romero github.com/jorgeaguirre150-source

15

YEARS INTERNATIONAL IT

5+

YEARS CLOUD
ARCHITECTURE

-30%

AWS COST VIA IAC

3

REGIONS
EMEA·APAC·NAFTA

Professional Summary

Senior Cloud Infrastructure Architect with **5+ years** designing enterprise multi-cloud platforms (Azure, AWS), built on **15 years** of international IT delivery across the EU and US. Currently co-owning the **mAzure HUB** platform at **Mercedes-Benz** — a multi-region hub-and-spoke Azure environment providing on-prem connectivity, shared services and **AI / GPU workload enablement**. Specialised in **Terraform / Terragrunt**, multi-tenant Kubernetes (AKS / EKS, vCluster patterns), **GPU platform engineering** (NVIDIA NGC / CUDA, GPU Operator) and **agentic AI architecture** (LangChain, LlamaIndex, RAG). Delivered a **30% reduction in AWS infrastructure cost** through IaC adoption and module reuse.

Azure

AWS

Terraform / Terragrunt

Kubernetes · AKS / EKS

GitLab CI/CD

NVIDIA NGC / CUDA

GPU Operator

Multi-Tenant GPU

LangChain / LangGraph

LlamaIndex

Anthropic SDK · RAG

ExpressRoute · Hub-Spoke

Prometheus · Grafana

Professional Experience

Senior Cloud Infrastructure Architect · Mercedes-Benz (mHUB Team)

Dec 2022 – Present · Stuttgart (Remote)

Lead digital-transformation within a multi-cloud strategy — security, automation and operational efficiency.

- ▶ **Terraform & Terragrunt** — reusable modules (networking, firewalls, VMSS, DNS, spoke peering) adopted across MBM product teams.
- ▶ **Multi-tenant Kubernetes & GPU** — tenant-spoke isolation on AKS/EKS (vCluster-style); self-service GPU environments under enterprise compliance.
- ▶ **Minerva — NVIDIA AI Factory** — principal architect: GPU Operator, NVIDIA device plugin, per-tenant GPU scheduling, identity-bound secure runtime for agentic workloads.
- ▶ **GitLab CI/CD** — IaC validate/plan/deploy pipelines enforcing security & compliance baselines across all environments.
- ▶ **Multi-cloud delivery** — Azure (primary), AWS (production), GCP (exposure) across EMEA / APAC / NAFTA.

★ **Principal architect — Minerva:** Mercedes-Benz AI-factory pattern for GPU workload enablement on NVIDIA NGC/CUDA.

★ **–30% AWS infrastructure cost** through IaC adoption and module reuse.

★ **Co-owner of mAzure HUB** — multi-tenant shared services for spokes across EMEA / APAC / NAFTA.

★ Observability with **Grafana / Prometheus**; post-mortem & DevOps culture embedded day-to-day.

Cloud Infrastructure Architect · Centene (USA · Healthcare)

Oct 2020 – Nov 2022 · Remote

Fortune 500 US healthcare — led cloud adoption and CI/CD for mission-critical workloads (HIPAA).

- ▶ **Terraform & Ansible** — IaC provisioning + configuration; automated complex lifecycles and deployments.
- ▶ **Application delivery** — JFrog Artifactory for artifact traceability, Oracle WebLogic for API admin on HA Linux/Windows estates.
- ▶ **Critical production** — expert support across Azure workloads and VMware vSphere; multi-cloud (Azure, AWS, Oracle, GCP).

★ **Formal recommendation letters** from project leads — recognised professional excellence.

★ **2 years fully remote** across multiple US time zones and regulatory contexts.

★ **Recommendation:** “I recommend Jorge without reservation” — Carl W. Blocker Jr, IT Manager, Systems Engineering · Centene

[Read letter →](#)

International Projects

European Medicines Agency — Amsterdam

Freelance · Mar 2017 – Sep 2020 · Regulated
Pharma (GxP)

IT & AV infrastructure for a regulated EU pharma
agency during the post-Brexit relocation.

AB InBev APAC — Belgium

Freelance · Feb 2014 – Jan 2017

Advanced L2/L3 support & administration for 500+
corporate users at APAC HQ.

BT Group — London

Mar 2012 – Jan 2014

Enterprise videoconferencing & advanced
communication technology for a high-profile
delivery.

BBVA — Madrid

Jan 2007 – Feb 2012 · Banking

Foundational IT support: administration of
infrastructure for 500+ banking users.

Strategic AI — Minerva & Pharma Optima

Functional Design & Agentic Implementation

Minerva (Mercedes-Benz · Enterprise Agentic AI) | Pharma Optima (Independent R&D · Regulated AI ·
2024–Present)

ServiceNow Integration

AI agents for autonomous request creation &
incident triage — reason over CMDB, apply
routing rules, escalate to humans below a
confidence threshold.

Validated Document AI

Multi-agent RAG for clinical data extraction with
GxP verifiability: planner / retriever / extractor /
verifier; every decision logged with prompt +
model version + citations.

GPU Orchestration

Agent workloads on NVIDIA NGC/CUDA via
mAzure HUB — private endpoints, identity-
bound access, per-tenant isolation. One paved
road for open-weights & hosted endpoints.

Risk Monitoring

Autonomous signal detection aligned with ICH
E6 — protocol deviations, adverse-event signals
& data-integrity alerts feeding the GxP audit
trail.

Claude

GPT-4o

Llama 3

LangChain / LangGraph

LlamaIndex

Anthropic SDK

NVIDIA NGC / CUDA

GPU Operator

vCluster / Multi-Tenant K8s

Run:AI / Ray / Jupyter

Azure / Terraform

Langfuse / Phoenix

GxP

ICH E6 R2

GDPR

21 CFR Part 11

ISO 27001

HIPAA

Core Competencies

01 Architectural Design — multi-region hub-and-spoke across Azure, AWS, Oracle.

02 Multi-Tenant K8s & GPU — AKS/EKS at scale, vCluster patterns, GPU Operator, device plugin.

03 AI Factory & GPU Infra — self-service tenant clusters, bare-metal & cloud-GPU enablement.

04 Infrastructure as Code — Terraform, Terragrunt, GitLab CI/CD, reusable enterprise modules.

05 Hybrid Cloud & Networking — ExpressRoute, hub-spoke, NVA NAT, DNS extended, peering.

06 AI Solution Architecture — multi-agent systems, RAG pipelines, GPU workload enablement.

07 Agentic AI Stack — LangChain, LlamaIndex, Anthropic SDK, tool use, evals on regulated workloads.

08 Security & Regulated Env. — Automotive, Pharma (GxP), Healthcare (HIPAA), Banking.

09 Monitoring & Observability — Grafana, Prometheus, Langfuse / Phoenix on Kubernetes.

10 Enterprise Intl Delivery — EMEA / APAC / NAFTA governance & compliance.

Languages


Spanish Native



English C1



German A2 · improving



Education & Training

Higher Education — Network Security (CISCO) · 2005–2007

Higher Education — IT & Telecommunications · 2005–2008

AI / GenAI Architecture — LangChain, LlamaIndex, RAG, agentic systems.

AZ-305 — Azure Solutions Architect Expert · in progress.

Let's build your cloud & AI platform.

Available for senior architecture engagements — Cloud platforms, GPU/AI infrastructure & Agentic AI for regulated industries.

✉ aguirre_coslada@hotmail.com

☎ +34 687 172 900

🌐 [LinkedIn](#)